

What's New in Java 8

Optional Lab: Method Enhancements

Overview

In this lab you'll refactor some Java applications to take advantage of the various method enhancement features available in Java 8.

Source folders

Student project: `StudentJavaSE8`

Solution project: `SolutionJavaSE8`

Roadmap

There are 4 exercises in this lab, of which the last exercise is "if time permits". Here is a brief summary of the tasks you will perform in each exercise; more detailed instructions follow later:

1. Using method references to implement a functional interface
2. Defining default methods in an interface
3. Implementing a factory mechanism based on constructor references
4. (If Time Permits) Defining static methods in an interface

Exercise 1: Using method references to implement a functional interface

In the *Student* project, expand the `student.methods.functions` package. Take a look at the following Java files:

- `UnaryFunc.java`
This is a simple generic interface that has a single method named `perform()`. This method represents the idea of a unary function, i.e. a function that takes a single parameter, performs an operation on it, and returns a result.
- `Square.java` and `Cube.java`
These classes implement `UnaryFunc<Double>` to perform different kinds of operation. `Square` performs a mathematical square operation, and `Cube` performs a mathematical cube operation.
- `Main.java`
This is the entry point for the program. The `main()` method calls `doOp()` several times, passing in two parameters each time:
 - The 1st parameter is an operand, i.e. the value to do some operation on.
 - The 2nd parameter is a `UnaryFunc<T>` object, which represents the specific operation to perform.

Inside `doOp()`, we invoke `perform()` on the `UnaryFunc<T>` object. Depending on what type of object this actually is, this will either cause the square or the cube to be obtained.

In Java 7, you'd define a separate class to represent every different type of operation. We've just defined two so far (`Square` and `Cube`), but the situation would quickly get out of hand if you wanted to represent lots of different operations. Java 8 offers two simpler alternatives:

- Define a lambda expression to represent an inline implementation of the interface. This approach is appropriate if you have a one-off algorithm in just one place in your code.
- Use a method reference to point to an existing method in a class. This approach is appropriate if you want to reuse an algorithm in multiple places in your code.

Take a look at the following 2 classes, which provide several existing useful methods:

- `DistanceConverter.java`
This class has a couple of methods that convert distances from km to miles, and vice versa. Each method can fulfil the role of the `perform()` method in the `UnaryFunc<T>` interface, because they have the same signature.
- `WeightConverter.java`
This class has a couple of methods that convert weights from kg to pounds, and vice versa. Both of these methods have the same signature as the `perform()` method too.

Bearing all this in mind, extend `main()` as follows (use method references as appropriate):

- Call `doOp()` a couple of times, to convert a distance from km to miles and vice versa.
- Call `doOp()` a couple of times, to convert a weight from kg to pounds and vice versa.

Exercise 2: Defining default methods in an interface

Java 8 allows you to define default methods in an interface, as we discussed during the chapter. There are two main reasons for Java 8's support for default methods in an interface:

- It allows you to add new methods to existing interfaces, without breaking any existing implementation classes (which obviously don't implement these new methods yet).
- It allows you to define an interface with lots of methods, of which all-but-one have a default implementation. As we've said several times now, this is what we mean by "functional interfaces" in Java 8. The notion of functional interfaces is crucial in Java 8, because you can only use lambda expressions and/or method references in conjunction with functional interfaces.

With this in mind, add a default method to the `UnaryFunc<T>` interface. The method should be named `formatResult`, and should take a `T` as a parameter and return a `String`. The default method implementation should return a string with a format such as the following:

`"Result: parameter_value"`

Then modify `doOp()` in `Main.java`, so that it uses the `UnaryFunc<T>` object's `formatResult()` method to format the result. Display the string result on the console, as at present.

Note what you just did. The `doOp()` method can invoke `formatResult()` on the `UnaryFunc<T>` object, without having to modify every single implementation class to implement that method. Of course, the implementation classes *can* define a `formatResult()` method if desirable. For example, implement `formatResult()` in `Square` and `Cube` and re-run the program.

Exercise 3: Implementing a factory mechanism based on constructor references

Expand the `student.methods.factory` package. The package contains 4 Java files:

- `Logger.java`
This is a simple interface with a single method named `log()`.
- `BriefLogger.java`
This class implements the `Logger` interface to display brief log messages on the console. The idea is that client code can create a `BriefLogger` object at the start of a series of tasks, and then call `log()` several times to display a series of messages during these tasks.
- `VerboseLogger.java`
This class is similar to `BriefLogger`, except that it displays verbose messages (rather than brief messages) on the console.
- `Main.java`
This is the entry point for the program. The `main()` method calls `doSomeStuff()` several times. Each time it's called, `doSomeStuff()` creates a `BriefLogger` or a `VerboseLogger` (depending on the string parameter), and writes a series of log messages to the console. Run the program to verify you understand how it works.

`doSomeStuff()` is quite clumsy at the moment. It relies on a string parameter to decide what type of logger to create. A better design would be for `doSomeStuff()` to receive a factory object as a parameter, and then use the factory object to create the appropriate type of logger. To implement this factory mechanism, follow these steps:

- Define a new interface named `LoggerProvider`. The interface should have a single method named `getLogger()`, which returns a `Logger` object.

- *[aside]*-----

In Java 7, the next step would be to define different implementations of `LoggerProvider` to create and return different types of logger. For example:

- You might define a class named `BriefLoggerProvider`, and implement `getLogger()` to create and return a `BriefLogger` instance.
- You might define a class named `VerboseLoggerProvider`, and implement `getLogger()` to create and return a `VerboseLogger` instance.

In Java 8, you can avoid having to define specific factory classes. Instead, you can use constructor references to refer to the existing constructors in the `BriefLogger` and `VerboseLogger` classes. These constructors effectively fulfil the role of the `getLogger()` method, to create and return a `BriefLogger` or `VerboseLogger` object respectively. The following page describes how to make use of constructor references.

- *[end aside]*-----

- Here's how to make use of constructor references:
 - In `Main.java`, change the signature of `doSomeStuff()` so that it receives a `LoggerProvider` object rather than a `String`.
 - Modify all the calls to `doSomeStuff()` so that instead of passing in a string, you pass in an appropriate constructor reference (i.e. `VerboseLogger::new` or `BriefLogger::new`). The compiler will treat the constructors as if they were implementations of the `getLogger()` method, to create and return the appropriate type of logger object.
 - Modify the code inside `doSomeStuff()` so that it uses the `LoggerProvider` to create the appropriate type of logger object.

When you've made all these changes, run the program again to ensure it works as before. There are several benefits to the code as it now stands:

- You've removed the messy and error-prone string tests in `doSomeStuff()`.
- The code is more extensible. If you add new types of logger in the future, you don't need to worry about adding another branch to the string tests.
- You've avoided the overhead of traditional factory mechanisms in Java, which rely on a separate factory class for each kind of object that needs to be created (i.e. you didn't have to define classes such as `BriefLoggerProvider` and `VerboseLoggerProvider`). This can be a huge advantage in large-scale applications.

Exercise 4 (If Time Permits): Defining static methods in an interface

In this exercise you'll implement the Null Object pattern in the logger application. For more information about this pattern, see:

- http://en.wikipedia.org/wiki/Null_Object_pattern

To understand the *raison d'être* of the Null Object pattern, follow these steps:

- In `main()`, add another call to `doSomeStuff()`, passing in `null` for the `LoggerProvider` parameter (e.g. to indicate you don't want any logging this time).
- Run the program. When the program reaches the point where you pass `null` into `doSomeStuff()`, it crashes when it tries to call `getLogger()` on the `null` reference.

To implement the Null Object pattern in the program, you can take advantage of Java 8's support for static methods in interfaces. Follow these steps:

- In the `LoggerProvider` interface, add a static method named `getNullLogger()`. Implement the interface so that it returns a `Logger` object with a do-nothing implementation of the `log()` method (you can use a lambda expression to implement the do-nothing version of the `log()` method, if you like).
- Modify `doSomeStuff()` so that it checks if the incoming `LoggerProvider` parameter is `null`, and either creates a real logger or a "Null Object" logger as appropriate. Whichever type of logger is created, all the subsequent calls to `log()` are safe; they will either display brief messages, verbose messages, or no messages at all (in the case where a "Null Object" logger is being used).

Run the program. Verify the program displays (or doesn't display) log messages correctly.